PORTFOLIO CASE STUDY

AI Onboarding Assistant

Designing a production GenAI platform for a global financial services organisation — and the adoption strategy that made the whole system actually work.

Client name and internal system names redacted per NDA.

First production GenAI platform	RAG · Human-in-the-loop	352 users · NPS 41	2023 – 2024

Lavine Dukhan · linkedin.com/in/lavine-dukhan-ux

Impact.

352	41	251	$999K
users onboarded	Net Promoter Score	customer projects	efficiency savings

878.5 business days saved — equivalent to 3.66 full-time employees.

Publicly reported figures — 2024.

The product improved because people used it thoughtfully. Getting people to use it thoughtfully was a design problem.

Too much documentation. Not enough answers.

Customer-facing agents were navigating vast, constantly updated product documentation to answer complex questions in real time, with customers waiting. The cost was measured in delayed implementations and lost revenue.

> *"Historically, customers struggled with 'too much' documentation yet paradoxically needed more depth on specific topics. Customer-facing employees struggled too: scoping projects, explaining every element of a product, resolving complex issues that span multiple teams."*

The consequence:
Delayed customer projects · Inconsistent scoping · Hidden drag on time-to-revenue

THE SOLUTION

AI Onboarding Assistant

The organisation's first production GenAI platform — built entirely within its own cloud infrastructure.

- RAG architecture — ingests product documentation
- Generates cited answers in real time
- Human-in-the-loop review pipeline
- Fine-tuning from agent feedback
- Admin governance portal

Joined mid-flight.
Left as AI adoption lead.

PHASE 1 — ADMIN PORTAL

Owned the human-in-the-loop review interface where subject matter experts reviewed AI answers, approved or edited them, and fed quality data back into fine-tuning.

PHASE 2 — CHAT FEATURES

Expanded to own the citations sidecar and the home page / quick chat features, both of which emerged directly from office hours sessions I was running with beta agents.

PHASE 3 — AI ADOPTION LEAD

Role evolved beyond design into AI literacy — running training sessions and pre-release previews using agents' real work scenarios. Not tool training. Behaviour change.

Throughout all phases: cross-product design consistency — ensuring what was designed matched what was built, and the experience felt coherent across assistant and admin portal.

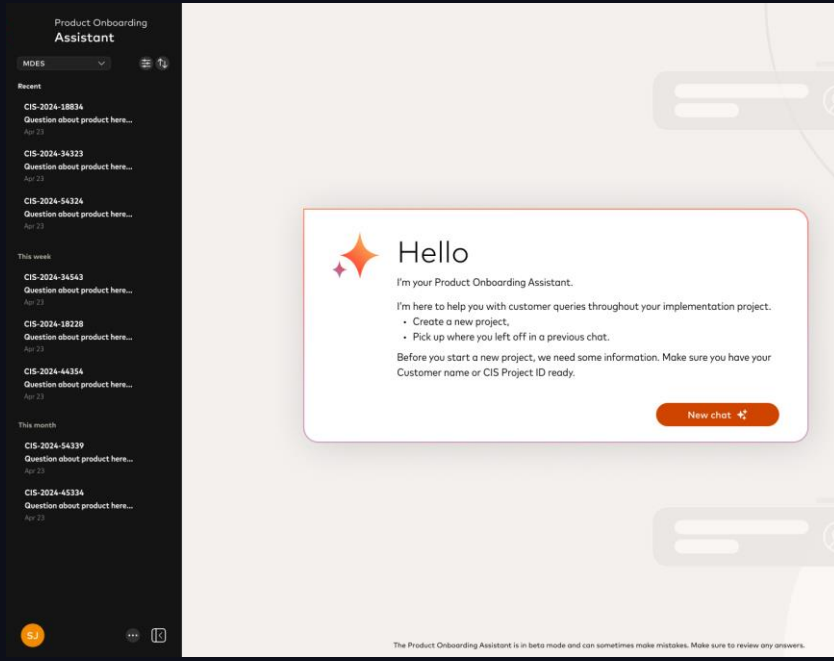

The first thing I did wasn't design. It was listen.

I introduced a regular office hours format with agents selected as testers. Not formal usability sessions, open conversations. A safe space for agents to share how they were actually experiencing the tool. This became a great Research Engine.

The insights from those sessions really helped us focus on what mattered to end users and it was a huge encouragement to the development team when they heard real-world stories from the people using what they were building.

These sessions became the primary intelligence source for every design decision I made for the rest of the project.

I regularly reported what agents were saying and tied it back to items already in the backlog proving the design strategy was on the right track before a single pixel shipped.

WHAT CAME OUT OF IT

Citations sidecar

Agents were verifying answers themselves, so I made that faster.

Home page design

Some agents were hesitant. Surfacing peer questions closed the gap.

Training content

Agents copying to Word. That was a literacy gap, not a UX problem.

Backlog validation

Tied every insight back to existing items, validating the strategy.

The instinct was to design the verification away. I made the opposite call.

Through office hours, I noticed agents receiving generated answers and then going to verify them against source documents before sending to customers. In a regulated financial environment, where product rules change frequently, this was entirely rational behaviour.

I didn't want to take away their ability to verify. I wanted to make verifying faster and more precise.

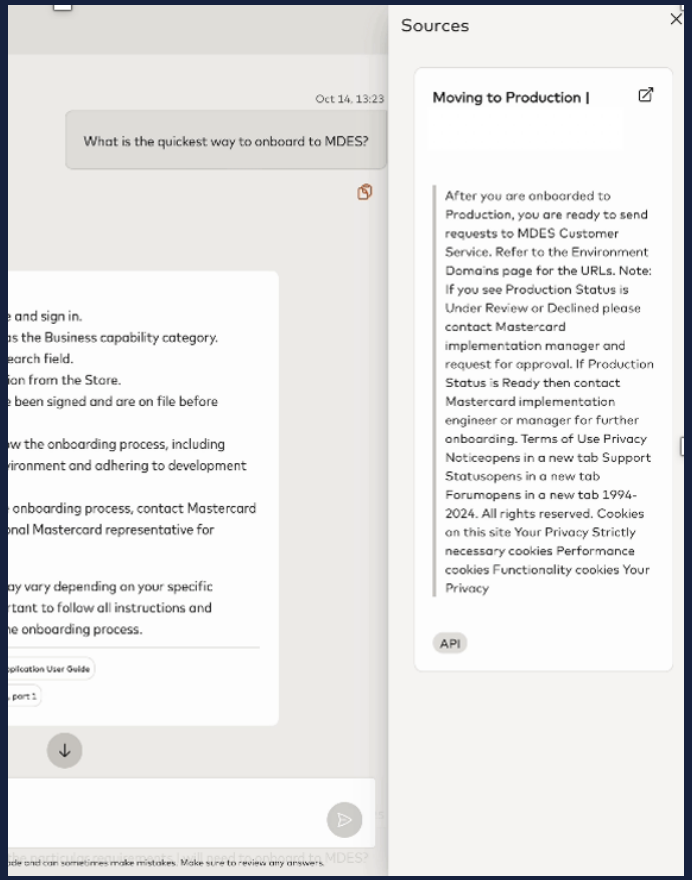

WHAT WE DIDN'T DO

- × Hide source info to seem more authoritative
- × Force agents to trust answers without routes to check
- × Add a separate verification workflow

WHAT WE BUILT INSTEAD

- ✓ Exact text snippets from source docs, inline
- ✓ Document name and metadata at a glance
- ✓ Dismissible: present when needed, invisible when not

Closing the system loop.

THE INSIGHT FROM OFFICE HOURS

Agents were copying AI answers into Word documents and reformatting them as tables or bullet points before sending to customers.

This wasn't a UX problem. It was a literacy gap. They didn't know they could ask the tool to format the answer in the first place.

PROMPT ENGINEERING

Every session used real agent workflows. 'Here is how you would have asked that question differently last Tuesday.' Specificity landed every time. Generic advice didn't.

FEEDBACK MATTERS

Agents who understood that their likes, dislikes, and edits fed the fine-tuning pipeline were significantly more likely to use feedback mechanisms consistently.

EDITING IS SAFE

Agents were nervous about editing AI answers. Explaining the human-in-the-loop review — every edited answer reviewed by an SME before retraining — removed that anxiety.

The result: usage spiked after every training session. The feedback loop that made the model smarter — likes, dislikes, edits — saw consistent increases following each cohort.

Some agents used the tool freely, others were hesitant. The design closed the gap.

The home page surfaced trending questions from other agents, filtered by the viewer's specific projects and products. The rationale was three layers deep; each one addressing a different dimension of the adoption gap.

01 CONFIDENCE BY PROXY

Agents who weren't ready to ask questions themselves could observe how others were using the tool and what useful answers looked like before committing.

02 PROMOTING CURIOSITY

Agents would often play around with the chat to see what kind of answers it gave and what they could do with it.

03 IMMEDIATE RELEVANCE

Filtered by the agent's own projects and products. A tailored tool for their specific work not a generic chatbot.

Quick chat was a sandbox; a space to see how this tool works for your products, your clients, your questions. A space to get comfortable before the stakes were real.

Assistant: V1

FAMILIAR PATTERNS

The feedback loop

The *feedback bar*, located below POA Q&A Pair, gives our users the opportunity to **rate the answers** and help improve their experience.

These ratings are then sent alongside the answers for a second level admin review, which is ultimately used to fine tune the model.

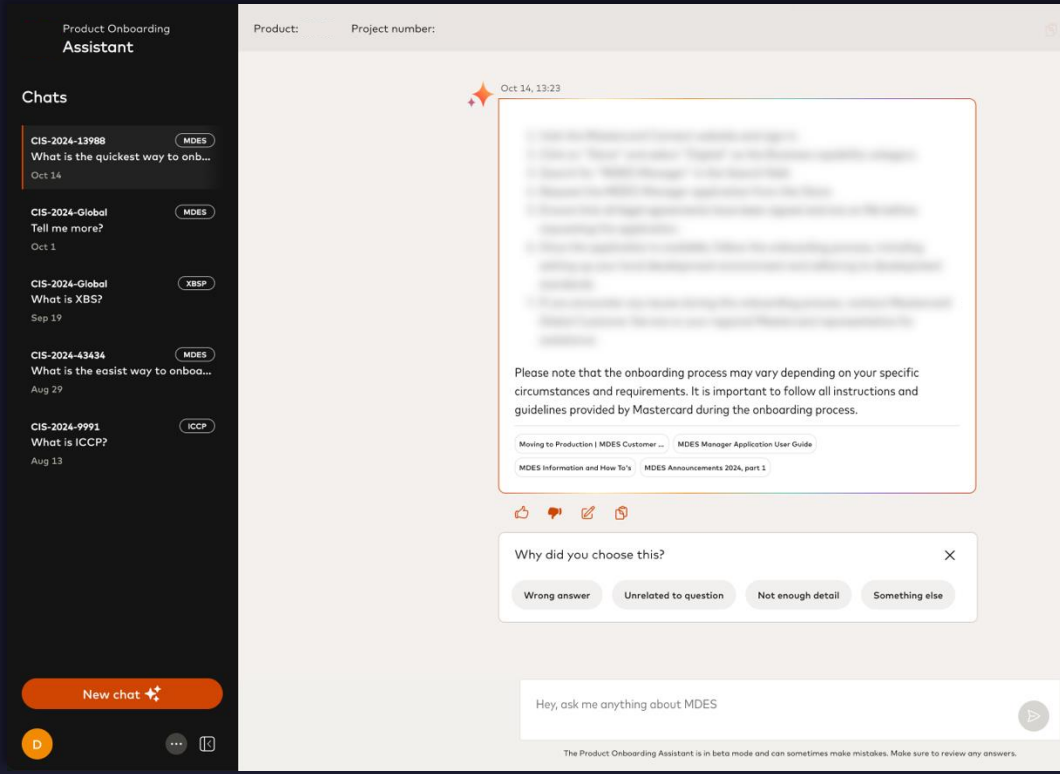

Agents saw what was coming — because they'd asked for it.

As new features were developed, I ran preview sessions with agents before each release showing them what was coming, and explaining how it would work so they knew exactly what would hitting their screens next.

AGENTS BECAME MORE ENGAGED

Conversations became more detailed, with agents describing real problems that they were facing. The tone shifted from them feeling like they were being judged on how they used the tool to them helping shape the tool

ZERO COLD-START

Agents arrived at each launch already familiar with what was coming and why. Adoption readiness replaced cold adoption friction.

DEV TEAM MOTIVATION

The development team attended these sessions. Hearing real user stories — not metrics — was a significant motivator and a visible proof point.

Agents could see that what they'd said had become what was being built. Closing the feedback loop visibly was as important as closing it technically.

The product improved because people used it thoughtfully.

Getting people to use it thoughtfully was a design problem.

THE ASSISTANT

Agent-facing chat. Questions, cited answers, feedback (like / dislike / edit), project-specific chat history.

THE ADMIN PORTAL

Human-in-the-loop curation. SMEs reviewed interactions, approved or corrected, prepared data for fine-tuning.

THE FEEDBACK PIPELINE

Every like, dislike, and edit was a data point. The edit feature was later adopted by the underlying engineering platform.

THE HUMAN LAYER

Training, office hours, previews, not adjacent to the product build but part of it. Better-trained agents gave better feedback. Better feedback made the model better.

This is what distinguished this project from a standard enterprise software deployment.

What I'd do differently. And why.

AI adoption is a design problem

Shipping an AI tool is the beginning, not the end. The gap between a capable system and a used system is entirely a design and behaviour change problem. Features don't drive adoption confidence does.

Trust precedes utility

Agents who didn't trust the tool didn't use it fully. Without trust, the feedback loop the whole system depended on couldn't function. The citations sidecar wasn't a feature it was the foundation.

Office hours — from day one

The insights from those conversations were more valuable than any formal research session. The earlier they arrive, the more they shape rather than retrofit.

Real examples are irreplaceable

Using the scenarios that came up in the office hours for the trainings helped make the learnings more tangible and gave the agents reassurance that their voices are being heard and made them more willing to attend office hours and training sessions

Appendix

Due to NDA the following examples have been altered to not reveal any company sensitive information

Assistant: Home & Quick chat

Building confidence through exploration and experiential learning

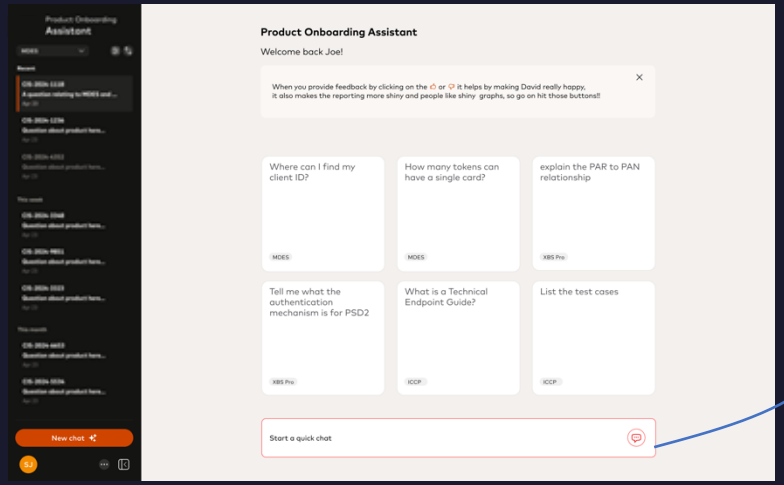

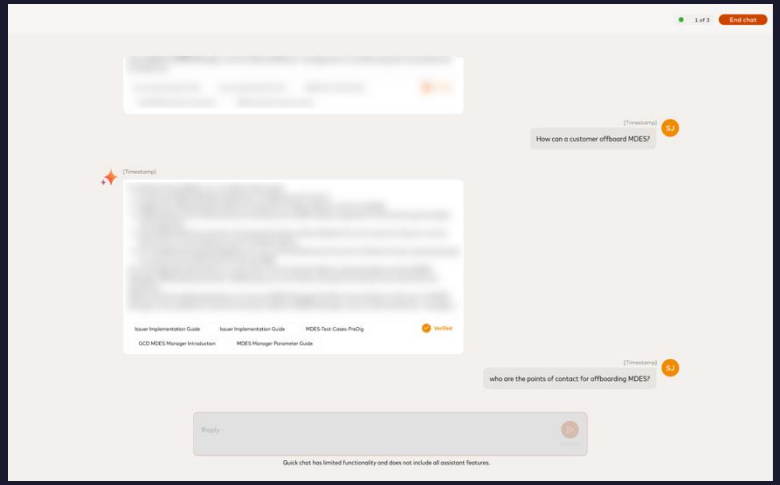

Home page

Agents see questions being asked for their products and are able to view the answers

- Targeted relevant content
- Learning from other agents questions
- Awareness around trending topics

Quick chat

Allows the agents to explore the product information and tool functionality by asking a quick question

- Safe experimentation area for agents to try new features
- Delve deeper into questions from other agents
- Enables agents to become more confident with AI tools